


# Applied Text Mining in Python

*Introduction to Text Mining*

# Text is Everywhere!



# Text data is growing fast!

- **Data continues to grow exponentially**
  - Estimated to be 2.5 Exabytes (2.5 million TB) a day
  - Grow to 40 Zettabytes (40 billion TB) by 2020 (50-times that of 2010)
- **Approximately 80% of all data is estimated to be unstructured, text-rich data**
  - >40 million articles (5 million in English) in Wikipedia
  - >4.5 billion Web pages
  - >500 million tweets a day, 200 billion a year
  - >1.5 trillion queries / searches on Google a year



# Data hidden in plain sight

Social  
network

Author

Description

Location

Tweet

- Topic
- Sentiment

Time

Popularity

The image shows a screenshot of the UN Spokesperson's Twitter profile. The profile header includes the UN logo, the name 'UN Spokesperson' with a verified badge, and the handle '@UN\_Spokesperson'. The bio states: 'Official Twitter account of the Office of the Spokesperson for United Nations Secretary-General Ban Ki-moon.' The location is 'New York, USA', the website is 'un.org/sg/spokesperso...', and it was joined in 'May 2010'. There is a button to 'Tweet to UN Spokesperson' and a link to '3,008 Photos and videos'.

The 'Tweets' tab is selected, showing three tweets. The first tweet is from 3 hours ago, discussing security challenges on the Korean Peninsula. The second tweet is from 17 hours ago, discussing ethics at the NY Society for Ethical Culture. The third tweet is from 23 hours ago, discussing the ban on Ambassador Joseph V. Reed.

Annotations highlight specific data points:

- Social network:** Points to the UN logo.
- Author:** Points to the profile name 'UN Spokesperson'.
- Description:** Points to the bio text.
- Location:** Points to the location 'New York, USA'.
- Tweet:** Points to the text of the second tweet.
- Time:** Points to the time '17h' of the second tweet.
- Popularity:** Points to the engagement metrics (6 replies, 13 retweets, 27 likes) of the second tweet.

Tweets	Following	Followers	Likes	Lists
14.6K	994	391K	49	3



# So, what can be done with text?

- Parse text
- Find / Identify / Extract relevant information from text
- Classify text documents
- Search for relevant text documents
- Sentiment analysis
- Topic modeling
- ...

# Applied Text Mining in Python

*Handling Text in Python*

# Primitive constructs in Text

- Sentences / input strings
- Words or Tokens
- Characters
- Document, larger files

And their properties ...

# Let's try it out!

```
>>> text1 = "Ethics are built right into the ideals and objectives  
of the United Nations "  
>>> len(text1)  
76  
>>> text2 = text1.split(' ')  
>>> len(text2)  
13  
>>> text2  
['Ethics', 'are', 'built', 'right', 'into', 'the', 'ideals', 'and',  
'objectives', 'of', 'the', 'United', 'Nations', '']
```



# Finding specific words

- **Long words:** Words that are most than 3 letters long

```
>>> [w for w in text2 if len(w) > 3]
['Ethics', 'built', 'right', 'into', 'ideals', 'objectives', 'United',
'Nations']
```

- **Capitalized words**

```
>>> [w for w in text2 if w.istitle()]
['Ethics', 'United', 'Nations']
```

- **Words that end with s**

```
>>> [w for w in text2 if w.endswith('s')]
['Ethics', 'ideals', 'objectives', 'Nations']
```

# Finding unique words: using set()

```
>>> text3 = 'To be or not to be'
>>> text4 = text3.split(' ')
>>> len(text4)
6
>>> len(set(text4))
5
>>> set(text4)
set(['not', 'To', 'or', 'to', 'be'])
>>> len(set([w.lower() for w in text4]))
4
>>> set([w.lower() for w in text4])
set(['not', 'to', 'or', 'be'])
```

# Some word comparison functions ...

- `s.startswith(t)`
- `s.endswith(t)`
- `t in s`
- `s.isupper()`; `s.islower()`; `s.istitle()`
- `s.isalpha()`; `s.isdigit()`; `s.isalnum()`



# String Operations

- `s.lower(); s.upper(); s.titlecase()`
- `s.split(t)`
- `s.splitlines()`
- `s.join(t)`
- `s.strip(); s.rstrip()`
- `s.find(t); s.rfind(t)`
- `s.replace(u, v)`

# From words to characters

```
>>> text5 = 'ouagadougou'
>>> text6 = text5.split('ou')
>>> text6
['', 'agad', 'g', '']
>>> 'ou'.join(text6)
'ouagadougou'
```

```
>>> text5.split('')
Traceback (most recent call last):
  File "<stdin>", line 1, in
<module>
ValueError: empty separator
>>> list(text5)
['o', 'u', 'a', 'g', 'a', 'd',
'o', 'u', 'g', 'o', 'u']
>>> [c for c in text5]
['o', 'u', 'a', 'g', 'a', 'd',
'o', 'u', 'g', 'o', 'u']
```

# Cleaning Text

```
>>> text8 = '    A quick brown fox jumped over the lazy dog. '
>>> text8.split(' ')
['', '', '\t', 'A', 'quick', 'brown', 'fox', 'jumped', 'over',
'the', 'lazy', 'dog.', '']
>>> text9 = text8.strip()
>>> text9.split(' ')
['A', 'quick', 'brown', 'fox', 'jumped', 'over', 'the', 'lazy',
'dog.']
```



# Changing Text

- Find and replace

```
>>> text9
'A quick brown fox jumped over the lazy dog.'
>>> text9.find('o')
10
>>> text9.rfind('o')
40
>>> text9.replace('o', 'O')
'A quick brOwn fOx jumped Over the lazy dOg.'
```

# Handling Larger Texts

- **Reading files line by line**

```
>>> f = open('UNDHR.txt', 'r')
>>> f.readline()
'Universal Declaration of Human Rights\n'
```

- **Reading the full file**

```
>>> f.seek(0)
>>> text12 = f.read()
>>> len(text12)
10891
>>> text13 = text12.splitlines()
>>> len(text13)
158
>>> text13[0]
'Universal Declaration of Human Rights'
```

# File Operations

- `f = open(filename, mode)`
- `f.readline(); f.read(); f.read(n)`
- `for line in f: doSomething(line)`
- `f.seek(n)`
- `f.write(message)`
- `f.close()`
- `f.closed`



# Issues with reading text files

```
>>> f = open('UNDHR.txt', 'r')
>>> text14 = f.readline()
'Universal Declaration of Human Rights\n'
```

- **How do you remove the last newline character?**

```
>>> text14.rstrip()
'Universal Declaration of Human Rights'
```

- Works also for DOS newlines (^M) that shows up as '`\r`' or '`\r\n`'

# Take Home Concepts

- Handling text sentences
- Splitting sentences into words, words into characters
- Finding unique words
- Handling text from documents

# Applied Text Mining in Python

## *Regular Expressions*



# Processing Free-text

```
>>> text10 = '"Ethics are built right into the ideals and  
objectives of the United Nations" #UNSG @ NY Society for Ethical  
Culture bit.ly/2guVelr @UN @UN_Women'  
>>> text11 = text10.split(' ')  
>>> text11  
['"Ethics', 'are', 'built', 'right', 'into', 'the', 'ideals',  
'and', 'objectives', 'of', 'the', 'United', 'Nations"', '#UNSG',  
'@', 'NY', 'Society', 'for', 'Ethical', 'Culture', 'bit.ly/  
2guVelr', '@UN', '@UN_Women']
```

- **How do you find all Hashtags? Callouts?**

# Finding Specific Words

- **Hashtags**

```
>>> [w for w in text11 if w.startswith('#')]
['#UNSG']
```

- **Callouts**

```
>>> [w for w in text11 if w.startswith('@')]
['@', '@UN', '@UN_Women']
```



# Finding patterns with regular expressions

- Callouts are more than just tokens beginning with '@'

@UN\_Spokesperson

@katyperry

@coursera

- Match something after '@'

- Alphabets
- Numbers
- Special symbols like '\_'

@**[A-Za-z0-9\_]+**

# Let's try it out!


```
>>> text10 = '"Ethics are built right into the ideals and objectives of the  
United Nations" #UNSG @ NY Society for Ethical Culture bit.ly/2guVelr @UN  
@UN_Women'  
>>> text11 = text10.split(' ')  
>>> [w for w in text11 if w.startswith('@')]  
['@', '@UN', '@UN_Women']
```

## Import regular expressions first!

```
>>> import re  
>>> [w for w in text11 if re.search('@[A-Za-z0-9_]+', w)]  
['@UN', '@UN_Women']
```



# Parsing the callout regular expression



- starts with @
- followed by any alphabet (upper or lower case), digit, or underscore
- that repeats at least once, but any number of times

# Meta-characters: Character matches

- `.` : wildcard, matches a single character
- `^` : start of a string
- `$` : end of a string
- `[ ]` : matches one of the set of characters within `[ ]`
- `[a-z]` : matches one of the range of characters a, b, ..., z
- `[^abc]` : matches a character that is not a, b, or c
- `a|b` : matches either a or b, where a and b are strings
- `( )` : Scoping for operators
- `\` : Escape character for special characters (`\t`, `\n`, `\b`)

# Meta-characters: Character symbols

**\b : Matches word boundary**

**\d : Any digit, equivalent to [ 0–9 ]**

**\D : Any non-digit, equivalent to [ ^0–9 ]**

**\s : Any whitespace, equivalent to [ \t\n\r\f\v ]**

**\S : Any non-whitespace, equivalent to [ ^ \t\n\r\f\v ]**

**\w : Alphanumeric character, equivalent to [ a-zA-Z0–9\_ ]**

**\W : Non-alphanumeric, equivalent to [ ^a-zA-Z0–9\_ ]**

# Meta-characters: Repetitions

- \* : matches zero or more occurrences**
- + : matches one or more occurrences**
- ? : matches zero or one occurrences**
- {n} : exactly n repetitions,  $n \geq 0$**
- {n,} : at least n repetitions**
- {,n} : at most n repetitions**
- {m,n} : at least m and at most n repetitions**

# Recall the callout regular expression

```
>>> text10 = '"Ethics are built right into the ideals and  
objectives of the United Nations" #UNSG @ NY Society for Ethical  
Culture bit.ly/2guVelr @UN @UN_Women'  
>>> text11 = text10.split(' ')  
  
>>> [w for w in text11 if re.search('@[A-Za-z0-9_]+', w)]  
['@UN', '@UN_Women']  
  
>>> [w for w in text11 if re.search('@\w+', w)]  
['@UN', '@UN_Women']
```



# Let's look at some more examples!

- Finding specific characters

```
>>> text12 = 'ouagadougou'
```

```
>>> re.findall(r'[aeiou]', text12)
['o', 'u', 'a', 'a', 'o', 'u', 'o', 'u']
```

```
>>> re.findall(r'^[aeiou]', text12)
['g', 'd', 'g']
```

# Case study: Regular expression for Dates

- Date variations for 23<sup>rd</sup> October 2002

23-10-2002

23/10/2002

23/10/02

10/23/2002

23 Oct 2002

23 October 2002

Oct 23, 2002

October 23, 2002

`\d{2}[/-]\d{2}[/-]\d{4}`

# Regular Expression for Dates (contd.)

```
>>> dateStr = '23-10-2002\n23/10/2002\n23/10/02\n10/23/2002\n23 Oct 2002\n23  
October 2002\nOct 23, 2002\nOctober 23, 2002\n'
```

```
>>> re.findall(r'\d{2}[/-]\d{2}[/-]\d{4}', dateStr)  
['23-10-2002', '23/10/2002', '10/23/2002']
```

```
>>> re.findall(r'\d{2}[/-]\d{2}[/-]\d{2,4}', dateStr)  
['23-10-2002', '23/10/2002', '23/10/02', '10/23/2002']
```

```
>>> re.findall(r'\d{1,2}[/-]\d{1,2}[/-]\d{2,4}', dateStr)  
['23-10-2002', '23/10/2002', '23/10/02', '10/23/2002']
```

23-10-2002
23/10/2002
23/10/02
10/23/2002

# Regex for Dates (contd.)

23 Oct 2002  
23 October 2002  
Oct 23, 2002  
October 23, 2002

```
>>> re.findall(r'\d{2} (Jan|Feb|Mar|Apr|May|Jun|Jul|Aug|Sep|Oct|Nov|Dec) \d{4}', dateStr)
['Oct']
```

```
>>> re.findall(r'\d{2} (? :Jan|Feb|Mar|Apr|May|Jun|Jul|Aug|Sep|Oct|Nov|Dec) \d{4}', dateStr)
['23 Oct 2002']
```

```
>>> re.findall(r'\d{2} (? :Jan|Feb|Mar|Apr|May|Jun|Jul|Aug|Sep|Oct|Nov|Dec) [a-z]* \d{4}', dateStr)
['23 Oct 2002', '23 October 2002']
```

# Regex for Dates (contd.)

23 Oct 2002  
23 October 2002  
Oct 23, 2002  
October 23, 2002

```
>>> re.findall(r'(?:\d{2} )?(?:Jan|Feb|Mar|Apr|May|Jun|Jul|Aug|Sep|Oct|Nov|  
Dec)[a-z]* (?:\d{2}, )?\d{4}', dateStr)  
['23 Oct 2002', '23 October 2002', 'Oct 23, 2002', 'October 23, 2002']
```

```
>>> re.findall(r'(?:\d{1,2} )?(?:Jan|Feb|Mar|Apr|May|Jun|Jul|Aug|Sep|Oct|Nov|  
Dec)[a-z]* (?:\d{1,2}, )?\d{4}', dateStr)  
['23 Oct 2002', '23 October 2002', 'Oct 23, 2002', 'October 23, 2002']
```

# Take Home Concepts

- **What are regular expressions?**
- **Regular expression meta-characters**
- **Building a regular expression to identify dates**



# Applied Text Mining in Python

*Internationalization*

# World of Languages



# English and ASCII

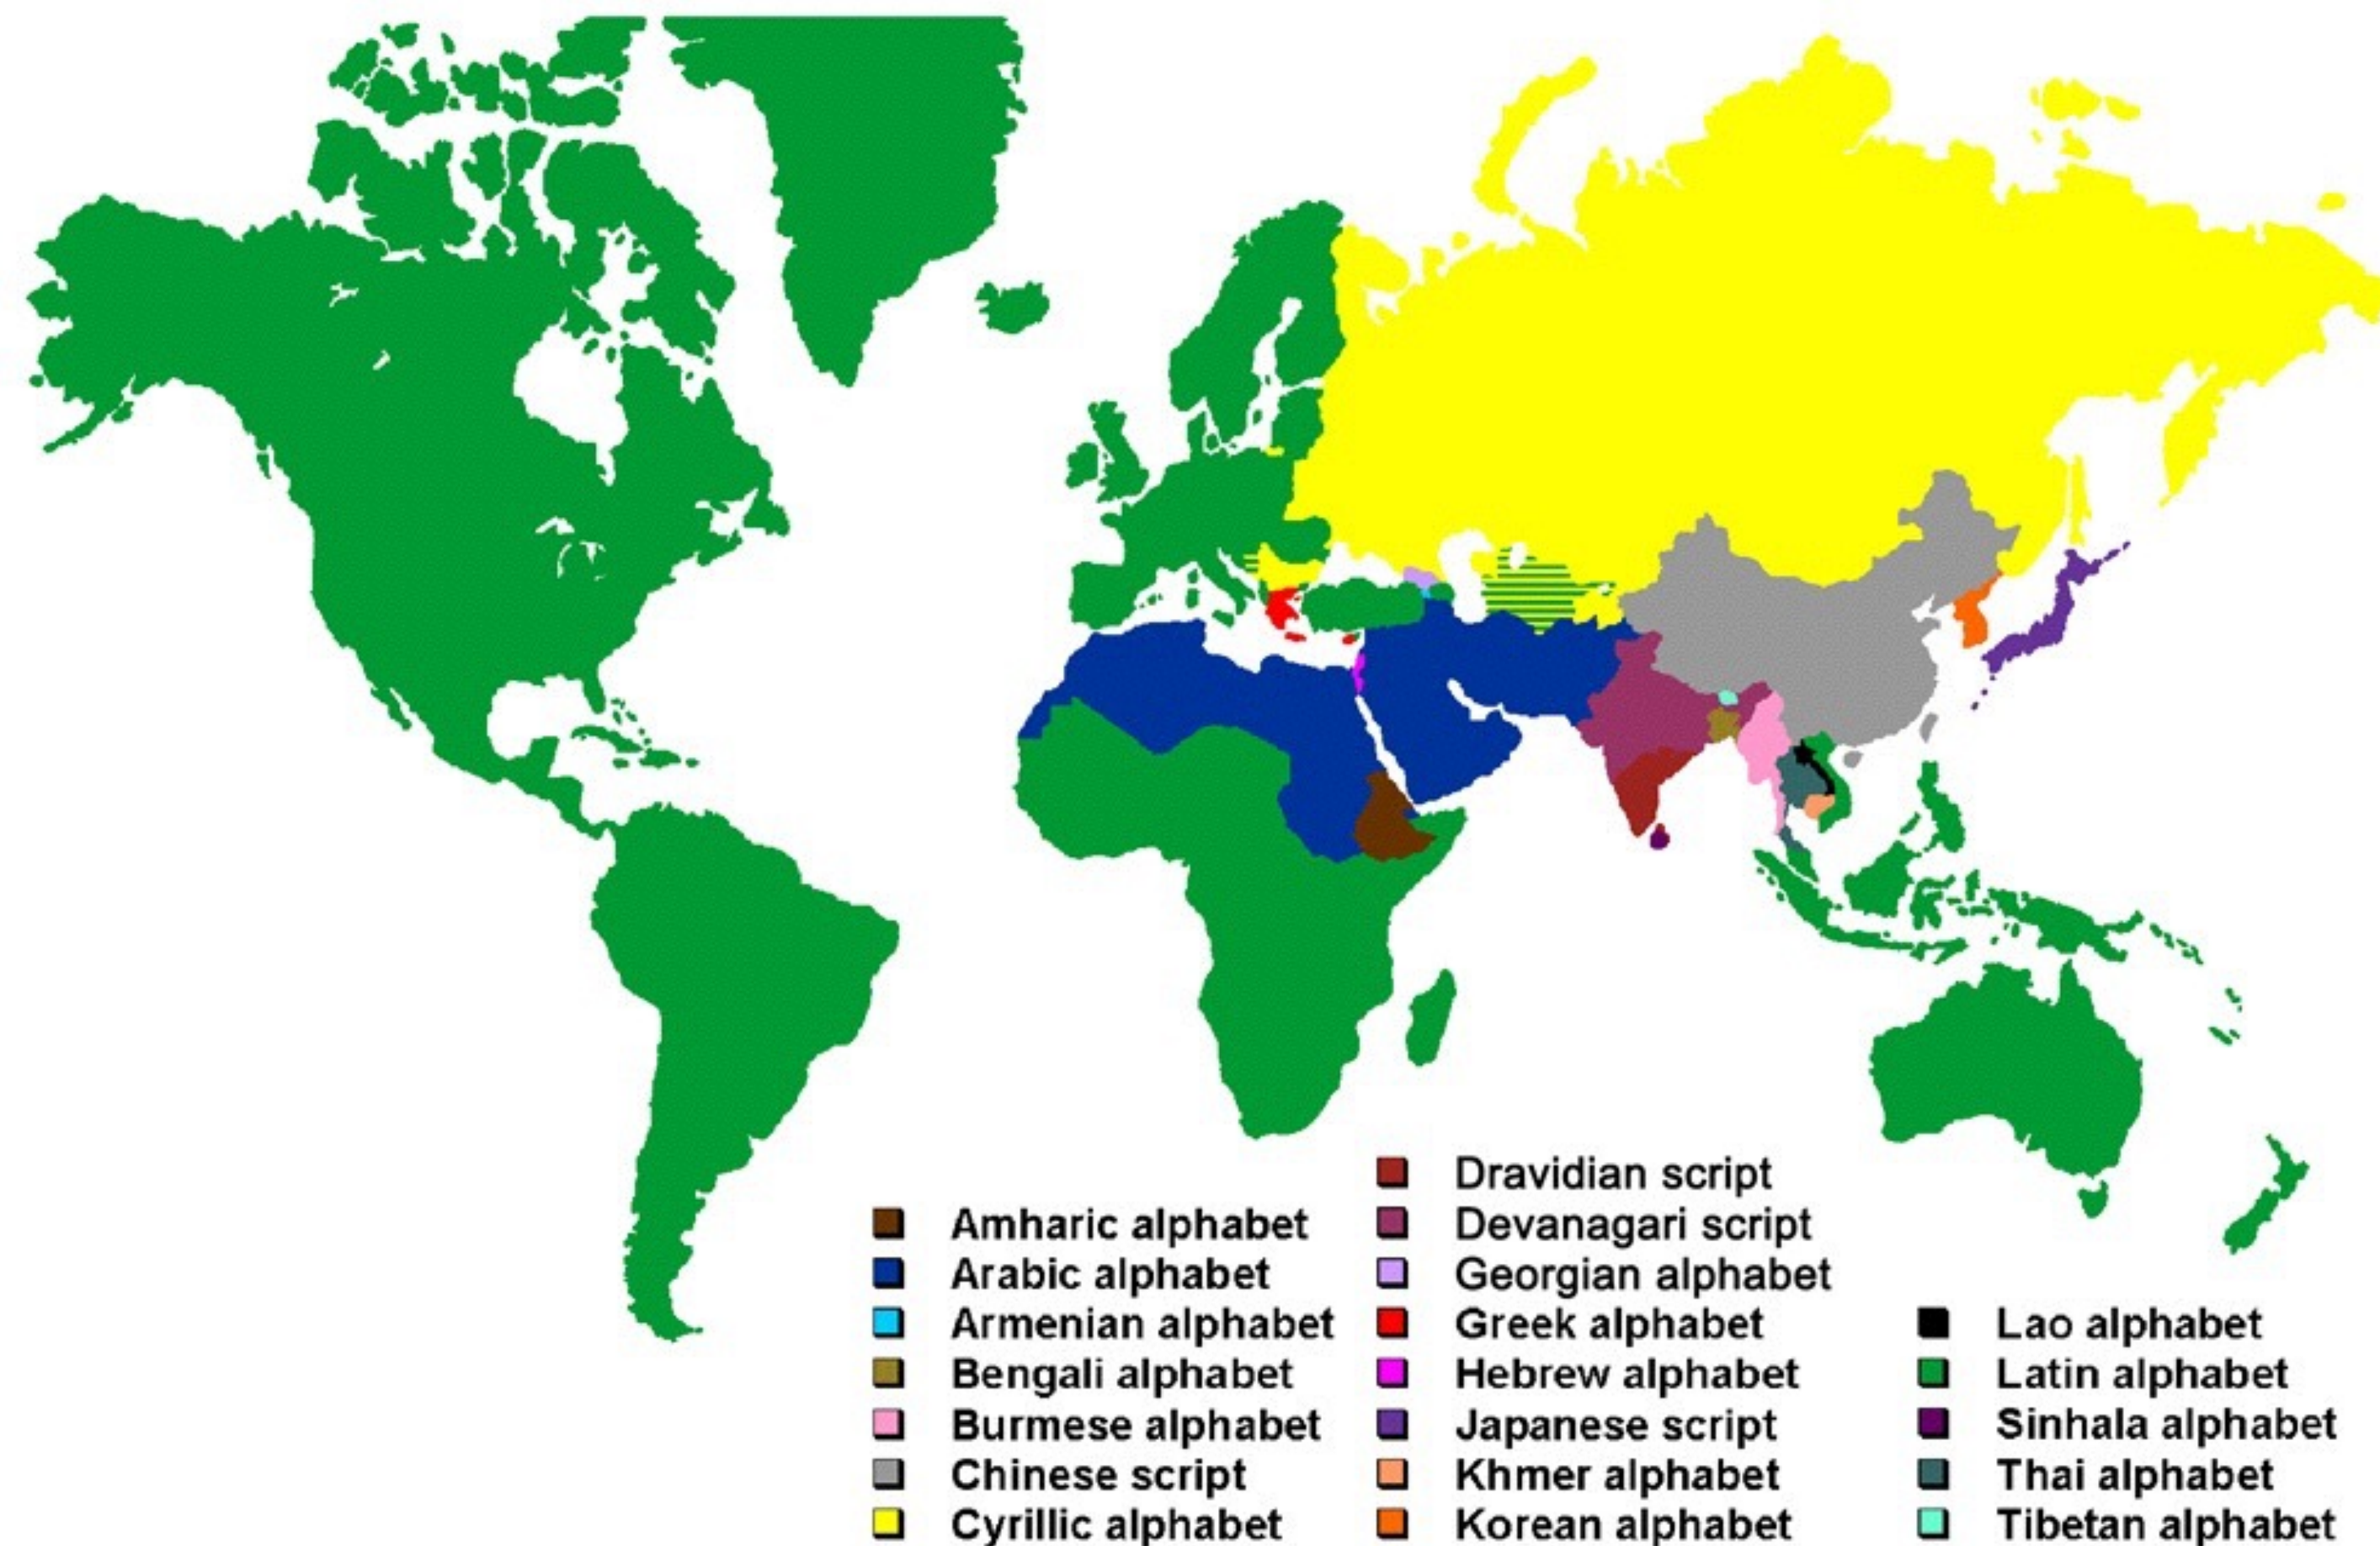
- **ASCII: American Standard Code for Information Interchange**
  - 7-bit character encoding standard: 128 valid codes
  - Range: 0x00 – 0x7F  $[(0000\ 0000)_2 \text{ to } (0111\ 1111)_2]$
  - Includes alphabets (upper and lower cases), digits, punctuations, common symbols, control characters
  - Worked (relatively) well for English typewriting

# Resume vs. Résumé

- **Diacritics**
  - résumé :: resume
  - naïve :: naive
  - café :: cafe
  - Québec
  - Zürich
  - Fédération Internationale de Football Association (FIFA)
- **International languages**
  - 基本上    सहायक    ασπασθ    универсальной
  - ♪    ♪    ♪
  - 😊    😞



# Written Scripts




- Latin: 36% (2.6B people)
- Chinese: 18% (1.3B)
- Devanagari: 14% (1B)
- Arabic: 14% (1B)
- Cyrillic: 4% (0.3B)
- Dravidian: 3.5% (0.25B)

# Other Character Encodings

- IBM EBCDIC
- Latin-I
- JIS: Japanese Industrial Standards
- CCCII: Chinese Character Code for Information Interchange
- EUC: Extended Unix Code
- Numerous other national standards
  
- Unicode and UTF-8



## Share of web pages with different encodings



# Unicode

- Industry standard for encoding and representing text
- Over 128,000 characters from 130+ scripts and symbol sets
- Can be implemented by different character endings
  - UTF-8: One byte to up to four bytes
  - UTF-16: One or two 16-bit code units
  - UTF-32: One 32-bit code unit

# UTF-8

- **Unicode Transformational Format – 8-bits**
- **Variable length encoding: One to four bytes**
- **Backward compatible with ASCII**
  - One byte codes same as ASCII
- **Dominant character encoding for the Web**
- **How to handle in Python?**
  - Default in Python 3
  - In Python 2:  
`# -*- coding: utf-8 -*-`

# Let's see an example: Résumé

## Python 3

```
>>> text1="Résumé"
>>> len(text1)
6
>>> text1
'Resumé'

>>> [c for c in text1]
['R', 'é', 's', 'u', 'm', 'é']
```

## Python 2

```
>>> text1="Résumé"
>>> len(text1)
8
>>> text1
'R\xc3\xa9sum\xc3\xa9'

>>> [c for c in text1]
['R', '\xc3', '\xa9', 's', 'u', 'm', '\xc3', '\xa9']

>>> text2=u'Resumé'
>>> len(text2)
6
>>> text2
u'R\xe9sum\xe9'
>>> [c for c in text2]
[u'R', u'\xe9', u's', u'u', u'm', u'\xe9']
```

# Take Home Concepts

- Diversity in Text
- ASCII and other character encodings
- Handling text in UTF-8